Hybrid versus standard tensor trains under time evolution: benchmark of the explicit solver with SVD truncation

S. Pisoni
TN_Turbulence project — 2D passive scalar study
(Dated: October 18, 2026)

The hybrid tensor-train encoding is benchmarked against a standard quantized tensor train for the time evolution of a passive scalar, at four resolutions from 256^2 to 2048^2. Both encodings are advanced by the same explicit solver: exact MPO–MPS contraction followed by a two-sweep SVD truncation. At a fixed target accuracy the hybrid encoding reaches that accuracy with a factor 1.8–2.5 fewer parameters, and the factor is constant across the resolutions tested. The head size that achieves it holds half the bits of each axis, $N \simeq L$, corresponding to a $2^{L/2} \times 2^{L/2}$ coarse grid represented densely and the finer scales carried by the tail. The error made in evolving tracks the error of merely representing the same field within a factor 1.5–25 for every encoding, resolution and budget measured. Two further flows are reported: a forced cellular flow, in which the rank plateau found in the decaying cascade persists in a statistically steady state, and a frozen DNS velocity snapshot, under which the rank also plateaus but the evolution error ceases to be monotone in the bond cap while the representation error remains so.

I. SETUP

The benchmark problem is a decaying passive scalar on a periodic square,

$$\partial_t \theta + \mathbf{u} \cdot \nabla \theta = \kappa \nabla^2 \theta, \qquad \nabla \cdot \mathbf{u} = 0, \tag{1}$$

with no forcing, so that the variance budget is

$$\frac{d}{dt}\frac{\langle \theta^2 \rangle}{2} = -\kappa \langle |\nabla \theta|^2 \rangle. \tag{2}$$

The velocity is the steady cellular flow obtained from the streamfunction $\psi = \sin x \, \sin y$, so $\mathbf{u} = (\sin x \cos y, \, -\cos x \sin y)$. It is constructed by applying the solver's own derivative MPOs to the encoded ψ, which makes it divergence-free to 7×10^{-13} at the discrete level and gives it bond dimension 4. The initial condition is a random-phase superposition of all modes with $|\mathbf{k}| \leq 2$, normalised to zero mean and unit variance.

Space is discretised with second-order centred stencils for both the transport and the diffusion terms, and the transport term is written in skew form, $\frac{1}{2}$(advective + conservative), which inherits the antisymmetry of the continuous operator and conserves $\langle \theta^2 \rangle$ to round-off in the absence of truncation. Time integration uses SSPRK3 with a step satisfying both the advective CFL condition and the two-dimensional diffusive limit $\Delta t \leq \Delta x^2 / 4\kappa$.

Resolution is varied at fixed $\ell_B / \Delta x = 2.88$ by scaling $\kappa \propto \Delta x^2$, anchored at $\kappa = 5 \times 10^{-3}$ for $L = 8$, so that each resolution resolves the same range of scales. All runs integrate to $T = 8$, which is past the time at which the truncation error saturates.

Both encodings are advanced by identical code. The standard tensor train is the case $N = 0$ of the hybrid encoder, verified to reproduce `tt.vector_to_mps(vec,2,2L)` bond for bond. Compression policy, contraction order, stencils, velocity, time step and dense reference are shared, so the head size is the only variable between the two.

II. HEAD SIZE: WHAT N CONTROLS

The field is encoded on $2L$ binary legs in Morton order, even legs carrying the row index and odd legs the column index, with level 0 the most significant bit. The head fuses the first N legs into a single dense tensor of physical dimension 2^N; the remaining $2L - N$ legs are ordinary binary cores.

N therefore selects how many of the *most significant* bits of each axis are held densely. The head holds $N/2$ bits per axis, which is a $2^{N/2} \times 2^{N/2}$ grid of coarse blocks, each block spanning $2^{L-N/2}$ cells on a side; the structure inside a block is carried by the tail. Table I gives the correspondence.

The useful range is $N \simeq L$, i.e. half the bits of each axis. The three sizes $N \in \{L-2, \, L, \, L+2\}$ bracket the optimum at every resolution measured, and the results below are reported on that bracket.

Two distinct normalisations of the head cost are worth separating, since they behave differently as L grows (Fig. 1). The head occupies $2^N \chi_0$ parameters. Relative to the compressed state this is a large and increasing share: at $L = 10$,

N bits/axis		blocks	cells per block		
	$N/2$	$2^{N/2} \times 2^{N/2}$	$L = 9$	$L = 10$	$L = 11$
4	2	4×4	128	256	512
6	3	8×8	64	128	256
8	4	16×16	32	64	128
10	5	32×32	16	32	64
12	6	64×64	8	16	32

TABLE I. Head size in terms of the grid it resolves densely. $N = 10$ at $L = 10$ holds 5 of the 10 bits of each axis, i.e. a 32×32 coarse grid, with the remaining 32×32 cells inside each block carried by the tail.

$N = 10$, $\chi = 32$ it is 74% of the stored parameters. Relative to the dense field it is $2^N \chi_0 / 4^L$, which for $N \simeq L$ is $\chi_0 / 2^L$ and therefore decreases with resolution: the same configuration stores a head worth 3.1% of the dense field, and the complete state worth 4.2%.

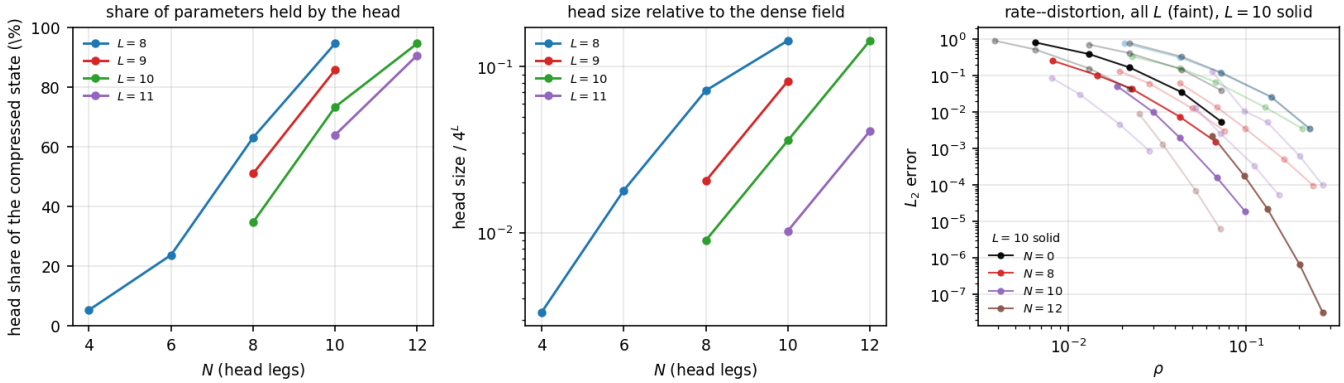


FIG. 1. Head cost under both normalisations, and the rate–distortion curves from which the optimum is read. Left: share of the compressed state held by the head. Centre: head size as a fraction of the dense field, which falls with L at fixed $N - L$. Right: L_2 error against compression ratio ρ for every run; $L = 10$ drawn solid.

III. RATE–DISTORTION

Bond caps are integers, so the two encodings do not land on identical parameter counts. Each encoding is therefore swept over bond caps and reported against the compression ratio $\rho = \text{parameters}/4^L$ it achieves. Table II lists every run. The reported error is the relative L_2 difference from a dense reference integrated with the same scheme, step and velocity, averaged over the final third of the run, where it has saturated.

Alongside each evolution error is the error of compressing the *exact* field at the same budget, which separates the cost of representing the solution from the additional cost of evolving in the representation.

IV. PARAMETERS REQUIRED AT FIXED ACCURACY

The quantity of practical interest is the budget required to reach a given accuracy. Table III reports the ratio $\rho_{\text{std}}/\rho_{\text{hybrid}}$, with the head size that achieves the smaller budget, obtained by log–log interpolation of $\rho(\text{error})$. Targets are listed only where both encodings span the target error.

The error ratio read at fixed ρ instead of fixed accuracy is considerably larger — at $L = 10$ it runs from 2.8 at $\rho = 0.01$ to 25 at $\rho = 0.05$ — and it grows with resolution. Both quantities describe the same pair of curves. The rate–distortion curves are steep and steepen with L, so a horizontal offset of a factor 2 in ρ appears as one to two decades along the error axis. The budget ratio is the quantity that answers how many parameters the head buys.

$L = 8$ (256^2)

N	χ	ρ	head %	L_2 (evolution)	L_2 (representation)
0	8	0.0221	–	7.96×10^{-1}	1.62×10^{-1}
0	12	0.0436	–	3.35×10^{-1}	6.92×10^{-2}
0	16	0.0729	–	1.22×10^{-1}	2.93×10^{-2}
0	24	0.1393	–	2.64×10^{-2}	7.30×10^{-3}
0	32	0.2291	–	3.57×10^{-3}	1.41×10^{-3}
4	8	0.0208	9	7.79×10^{-1}	1.62×10^{-1}
4	12	0.0423	7	3.22×10^{-1}	6.92×10^{-2}
4	16	0.0716	5	1.24×10^{-1}	2.93×10^{-2}
4	24	0.1380	3	2.66×10^{-2}	7.30×10^{-3}
4	32	0.2278	2	3.57×10^{-3}	1.41×10^{-3}
6	8	0.0228	34	3.44×10^{-1}	6.79×10^{-2}
6	12	0.0423	28	1.65×10^{-1}	2.98×10^{-2}
6	16	0.0677	23	6.76×10^{-2}	1.54×10^{-2}
6	24	0.1282	18	1.38×10^{-2}	4.34×10^{-3}
6	32	0.2083	15	3.58×10^{-3}	1.41×10^{-3}
8	8	0.0423	74	6.40×10^{-2}	1.57×10^{-2}
8	12	0.0687	68	1.39×10^{-2}	4.38×10^{-3}
8	16	0.0989	63	3.63×10^{-3}	1.59×10^{-3}
8	24	0.1634	57	5.21×10^{-4}	2.73×10^{-4}
8	32	0.2396	52	9.85×10^{-5}	5.14×10^{-5}
10	4	0.0648	97	1.30×10^{-1}	2.15×10^{-2}
10	6	0.0981	96	1.09×10^{-2}	4.78×10^{-3}
10	8	0.1321	95	5.32×10^{-3}	2.08×10^{-3}
10	12	0.2005	94	6.41×10^{-4}	3.46×10^{-4}
10	16	0.2708	92	1.00×10^{-4}	6.87×10^{-5}

$L = 9$ (512^2)

N	χ	ρ	head %	L_2 (evolution)	L_2 (representation)
0	12	0.0131	–	7.23×10^{-1}	1.36×10^{-1}
0	16	0.0221	–	4.20×10^{-1}	7.39×10^{-2}
0	24	0.0436	–	1.50×10^{-1}	2.33×10^{-2}
0	32	0.0729	–	3.95×10^{-2}	6.65×10^{-3}
8	12	0.0194	61	1.35×10^{-1}	1.85×10^{-2}
8	16	0.0286	55	6.20×10^{-2}	9.17×10^{-3}
8	24	0.0496	47	1.30×10^{-2}	2.30×10^{-3}
8	32	0.0755	41	3.13×10^{-3}	7.48×10^{-4}
10	12	0.0523	90	1.30×10^{-2}	1.93×10^{-3}
10	16	0.0716	87	2.69×10^{-3}	7.28×10^{-4}
10	24	0.1112	84	3.47×10^{-4}	1.19×10^{-4}
10	32	0.1536	81	5.32×10^{-5}	1.98×10^{-5}

$L = 10$ (1024^2)

N	χ	ρ	head %	L_2 (evolution)	L_2 (representation)
0	16	0.0065	–	8.28×10^{-1}	1.15×10^{-1}
0	24	0.0131	–	3.96×10^{-1}	4.51×10^{-2}
0	32	0.0221	–	1.70×10^{-1}	1.73×10^{-2}
0	48	0.0436	–	3.54×10^{-2}	3.55×10^{-3}
0	64	0.0729	–	5.46×10^{-3}	6.41×10^{-4}
8	16	0.0081	48	2.62×10^{-1}	2.54×10^{-2}
8	24	0.0146	40	1.03×10^{-1}	9.67×10^{-3}
8	32	0.0228	34	4.41×10^{-2}	4.10×10^{-3}
8	48	0.0423	28	7.60×10^{-3}	8.50×10^{-4}
8	64	0.0677	23	1.56×10^{-3}	2.16×10^{-4}
10	16	0.0189	83	5.19×10^{-2}	3.77×10^{-3}
10	24	0.0300	78	1.02×10^{-2}	9.06×10^{-4}
10	32	0.0423	74	2.03×10^{-3}	2.66×10^{-4}
10	48	0.0687	68	1.62×10^{-4}	3.19×10^{-5}
10	64	0.0990	63	1.92×10^{-5}	4.68×10^{-6}
12	16	0.0648	96	2.27×10^{-3}	3.20×10^{-4}
12	24	0.0981	96	1.80×10^{-4}	3.71×10^{-5}
12	32	0.1322	95	2.26×10^{-5}	6.34×10^{-6}
12	48	0.2005	94	6.80×10^{-7}	2.46×10^{-7}
12	64	0.2708	92	3.26×10^{-8}	1.06×10^{-8}

$L = 11$ (2048^2)

N	χ	ρ	head %	L_2 (evolution)	L_2 (representation)

target L_2	$L = 8$	$L = 9$	$L = 10$	$L = 11$
1×10^{-1}	1.4 (N=6)	2.3 (N=8)	1.9 (N=8)	–
6×10^{-2}	2.3 (N=8)	2.1 (N=8)	1.8 (N=8)	2.1 (N=10)
3×10^{-2}	2.5 (N=8)	–	2.1 (N=10)	–
1×10^{-2}	2.4 (N=8)	–	2.1 (N=10)	–
3×10^{-3}	–	–	–	–

TABLE III. Parameters saved at fixed accuracy, and the head size achieving it. The saving is 1.8–2.5 and shows no trend with resolution over 256^2–2048^2.

V. EVOLUTION AGAINST REPRESENTATION

Figure 2 collects the comparison across resolutions. The lower middle panel plots the evolution error against the representation error at the same budget for every run. All points lie between the 1× and 25× lines, with most below 5×, across four resolutions, five head sizes and five bond caps.

The practical consequence is that the ordering of encodings in N and the scaling in L can be established by compressing stored fields, which costs minutes, rather than by integrating them, which costs core-days. The representation measurement does not substitute for the evolution measurement when the absolute error level matters.

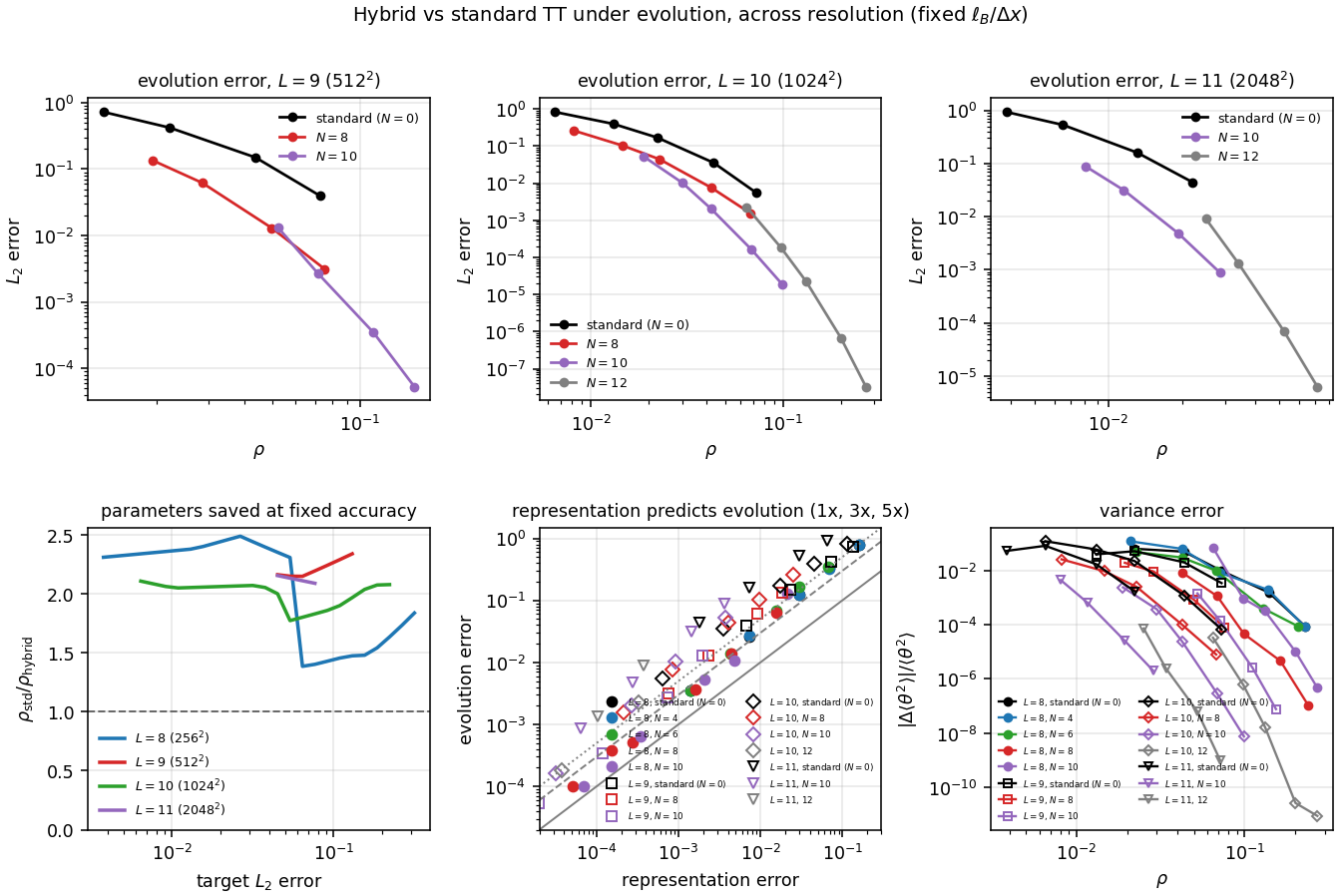


FIG. 2. Top: rate–distortion at $L = 9, 10, 11$. Bottom left: parameters saved at fixed accuracy. Bottom centre: evolution error against representation error at the same budget, with 1×, 3× and 5× guides. Bottom right: relative error in $\langle\theta^2\rangle$.

VI. COST

Table IV gives measured cost per step at $\chi = 32$. At equal bond cap the hybrid encoding is slightly cheaper than the standard one, since it carries $2L - N$ tail cores instead of $2L$, while the head contributes $O(2^N r^2 \chi)$ through the un-fusing path. At equal ρ the comparison is closer, because reaching a given ρ with a larger head permits a smaller χ.

L	N	χ	ρ	ms/step
8	0	32	0.2291	846
8	4	32	0.2278	920
8	6	32	0.2083	822
8	8	32	0.2396	535
9	0	32	0.0729	1507
9	8	32	0.0755	1419
9	10	32	0.1536	1518
10	0	32	0.0221	818
10	8	32	0.0228	940
10	10	32	0.0423	1068
10	12	32	0.1322	1989
11	0	32	0.0065	2257
11	10	32	0.0116	3348
11	12	32	0.0340	3596

TABLE IV. Measured cost per step at $\chi = 32$. Runs shared cluster nodes, so absolute values carry scheduling noise; the comparison within a row of constant L is the meaningful one.

All runs are single-threaded. On a 256-core node, four threads instead of one measured 29× slower per step, because a step consists of thousands of small contractions and SVDs, none large enough to amortise thread dispatch.

VII. FORCED STEADY STATE

The results above are obtained in a decaying cascade. To check that they are not an artefact of the transient, the cellular flow was rerun at 256^2 with a steady source $f = 0.8 \sin 4y$ to $T = 60$.

The forced wavenumber matters. The approach to a statistically steady state is governed by the diffusive time of the forced scale, $1/(\kappa k_f^2)$, which is 200 time units at $k_f = 1$; amplitudes from 0.02 to 0.05 were all still climbing at $t = 90$. At $k_f = 4$ that time is 12, and the variance plateaus by $t \simeq 40$, settling at 0.10, 0.22 and 0.51 for amplitudes 0.3, 0.5 and 0.8.

In the steady state the variance holds to within 1.3% beyond $t = 36$, and the intrinsic rank of the exact solution rises from 12 to a peak of 56 and settles at 41.9 ± 0.6 against a geometric bound of 256. The evolution error is monotone in χ for every head size, and the budget saving over the standard encoding is 1.4 at $L_2 = 10^{-2}$, achieved at $N = 6$ (Table V).

N	χ	ρ	head %	L_2 (evolution)	L_2 (representation)
0	16	0.0729	–	2.68×10^{-2}	3.14×10^{-3}
0	24	0.1393	–	2.86×10^{-3}	6.46×10^{-4}
0	32	0.2291	–	2.50×10^{-4}	1.05×10^{-4}
6	16	0.0677	23	1.22×10^{-2}	1.88×10^{-3}
6	24	0.1282	18	1.17×10^{-3}	4.28×10^{-4}
6	32	0.2083	15	2.50×10^{-4}	1.05×10^{-4}
8	16	0.0989	63	4.66×10^{-4}	2.21×10^{-4}
8	24	0.1634	57	5.68×10^{-5}	3.74×10^{-5}
8	32	0.2396	52	9.07×10^{-6}	6.35×10^{-6}
10	16	0.2708	92	1.32×10^{-5}	1.15×10^{-5}
10	24	0.4095	92	6.34×10^{-7}	5.20×10^{-7}
10	32	0.5521	91	5.14×10^{-8}	4.54×10^{-8}

TABLE V. Forced cellular flow, 256^2, $f = 0.8 \sin 4y$, averaged over the final third of the run.

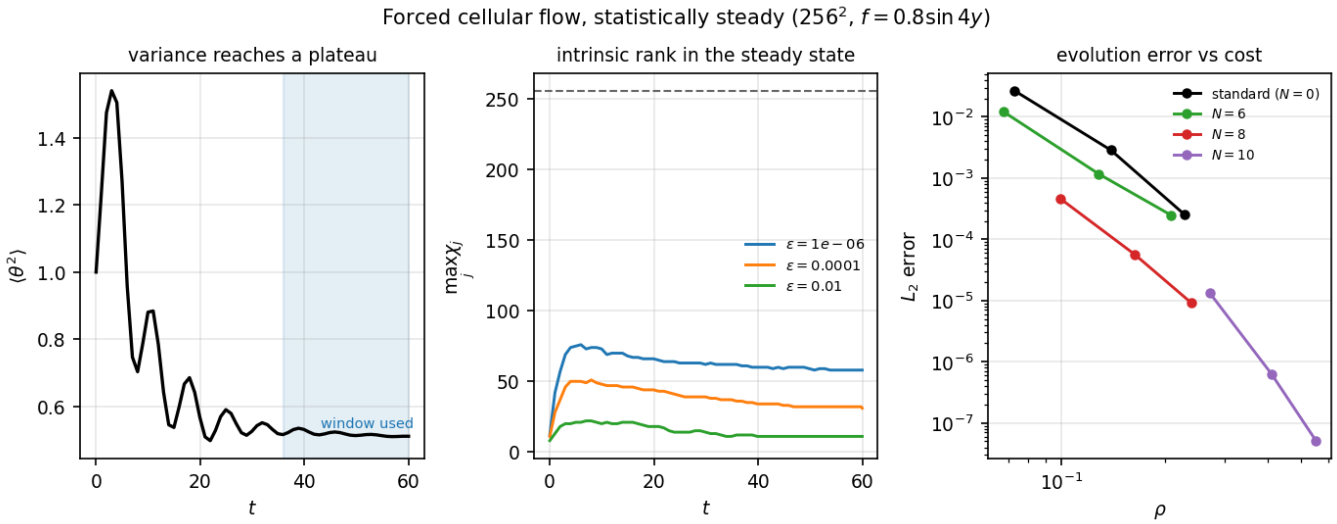


FIG. 3. Variance plateau, intrinsic rank in the steady state, and rate–distortion for the forced cellular flow.

VIII. A FROZEN DNS VELOCITY

The cellular flow is integrable and its scalar spectrum steep, so the benchmark was repeated with a velocity taken from the DNS dataset. The streamfunction was recovered from the stored v_x, v_y at 4096^2 by inverting the vorticity spectrally, coarsened to 512^2, and differentiated with the solver's own MPOs, leaving the velocity divergence-free to 2×10^{-12} discretely.

The streamfunction itself is close to full rank. At $\epsilon = 10^{-6}$ its bond saturates the geometric bound and the tensor train costs more parameters than the dense field, $\rho = 1.33$ at $L = 9$ and 2.32 at $L = 10$. It becomes compressible only at $\epsilon = 10^{-2}$, giving bond 21 at $\rho = 0.084$ for a 1.7% error in ψ. The runs therefore use a truncated velocity, carrying bond 34 against 4 for the cellular flow, which makes a step about five times more expensive. With $\kappa = 1.5 \times 10^{-2}$ the Batchelor scale is about three cells for the measured strain rate of 11, and $T = 2$ covers some 22 eddy turnover times.

The rank plateaus here as well, and faster: it is reached within a tenth of a time unit and then holds at 315 ± 1.1 at $\epsilon = 10^{-6}$ (61% of the geometric bound), 154 at 10^{-4} and 34 at 10^{-2}.

The evolution error, however, ceases to be monotone in the bond cap. At $N = 10$ the L_2 error runs 7.7×10^{-2}, 1.5×10^{-1}, 2.9×10^{-2}, 1.9×10^{-3} for $\chi = 16, 24, 32, 48$, and the variance error is non-monotone at $N = 8$, 10 and 12. The representation error at the same budgets is monotone in every case (Table VI), which locates the effect in the dynamics rather than in the encoding. The ratio of evolution to representation error widens from 1.5–25 for the cellular flow to 1.5–233 here.

Two consequences follow. The rate–distortion framing used above does not transfer directly, since a given (N, χ) cannot be ranked by its budget alone; and the budget saving can be quoted at only one accuracy target where the two encodings overlap, 2.1 at $L_2 = 10^{-1}$, achieved at $N = 8$.

IX. SCOPE OF THESE RESULTS

The measurements above are specific to the flow used. The cellular flow $\psi = \sin x \sin y$ has closed streamlines and is integrable, and its scalar spectrum is steep, $E_\theta(k) \sim k^{-4.9}$ over $3 \le k \le 20$, rather than the k^{-1} Batchelor form of a chaotic flow. Breaking the streamlines with $\psi = \sin(x + B \sin \omega t) \sin y$ raises the rank plateau from 75 to 84 at $L = 8$ and leaves the spectrum steep.

The size of the hybrid advantage depends on the field. On random-phase fields of prescribed spectral slope the same measurement gives 1.2 at $k^{-5/3}$ and 3.5 at k^{-5}, against 5–24 at fixed ρ for the cellular field. The difference is attributable to spatial coherence rather than to spectral slope or to alignment with the dyadic block boundaries: rigid shifts of the field by 2, 4, 8 and 13 cells leave the advantage unchanged.

Real turbulent scalar fields lie between the two cases. The DNS-velocity run above gives the first direct measurement on a realistic field: the advantage persists, at 2.1 in budget at $L_2 = 10^{-1}$, but the non-monotonicity of the evolution error means a single number characterises it less completely than in the cellular case. The representation measurements

N	χ	ρ	head %	L_2 (evolution)	L_2 (representation)
0	16	0.0221	–	3.06×10^{-1}	5.94×10^{-2}
0	24	0.0436	–	1.82×10^{-1}	2.97×10^{-2}
0	32	0.0729	–	1.25×10^{-1}	1.37×10^{-2}
0	48	0.1393	–	6.61×10^{-2}	5.03×10^{-3}
8	16	0.0286	55	1.51×10^{-1}	1.29×10^{-2}
8	24	0.0496	47	8.60×10^{-2}	6.46×10^{-3}
8	32	0.0755	41	8.44×10^{-2}	3.83×10^{-3}
8	48	0.1341	35	3.40×10^{-2}	1.85×10^{-3}
10	16	0.0716	87	7.00×10^{-2}	2.62×10^{-3}
10	24	0.1112	84	1.34×10^{-1}	1.26×10^{-3}
10	32	0.1536	81	2.68×10^{-2}	7.22×10^{-4}
10	48	0.2396	78	1.84×10^{-3}	3.26×10^{-4}
12	16	0.2552	98	9.06×10^{-2}	3.89×10^{-4}
12	24	0.3836	98	4.47×10^{-4}	1.54×10^{-4}
12	32	0.5130	97	1.14×10^{-4}	5.13×10^{-5}
12	48	0.7669	98	7.41×10^{-6}	5.04×10^{-6}

TABLE VI. Frozen DNS velocity, 512^2, ψ truncated at 10^{-2}, averaged over the final third of the run.

on DNS data reported in the paper remain the reference point for the representation question.

One further check bears on those figures. `tt.svd` defaults to `renormalize=True`, which rescales the kept singular values at every split. On `theta.00059` at 4096^2 the stored hybrid fields were found to equal the plain rank-truncated field times a single global scalar c to 1.2–3.5 $\times 10^{-14}$. The norm is preserved exactly, $E(1)$ shifts by exactly c^2 — +2.7% at $\rho = 0.014$, falling to +0.04% at $\rho = 0.64$ — and the gradient flatness ratio is 1.0000, so all scale-invariant statistics are unaffected.

X. CONFIGURATION

- Grids $2^L \times 2^L$ with $L = 8, 9, 10, 11$; $\kappa = 5 \times 10^{-3}$, 1.25×10^{-3}, 3.125×10^{-4}, 7.8125×10^{-5}.
- $\Delta t = 0.2\,\Delta x$ (CFL-limited at every resolution), $T = 8$, giving 1630, 3259, 6519 and 13038 steps.
- Head sizes $N \in \{0, 4, 6, 8, 10, 12\}$, bond caps $\chi \in \{4, \ldots, 64\}$ applied uniformly as `bd_first` = `bd_remaining`, with `renormalize` = `False` so that truncated norm is not restored.
- 69 runs in total; the $L = 10$ and $L = 11$ sets were run as Slurm arrays, one core per task.

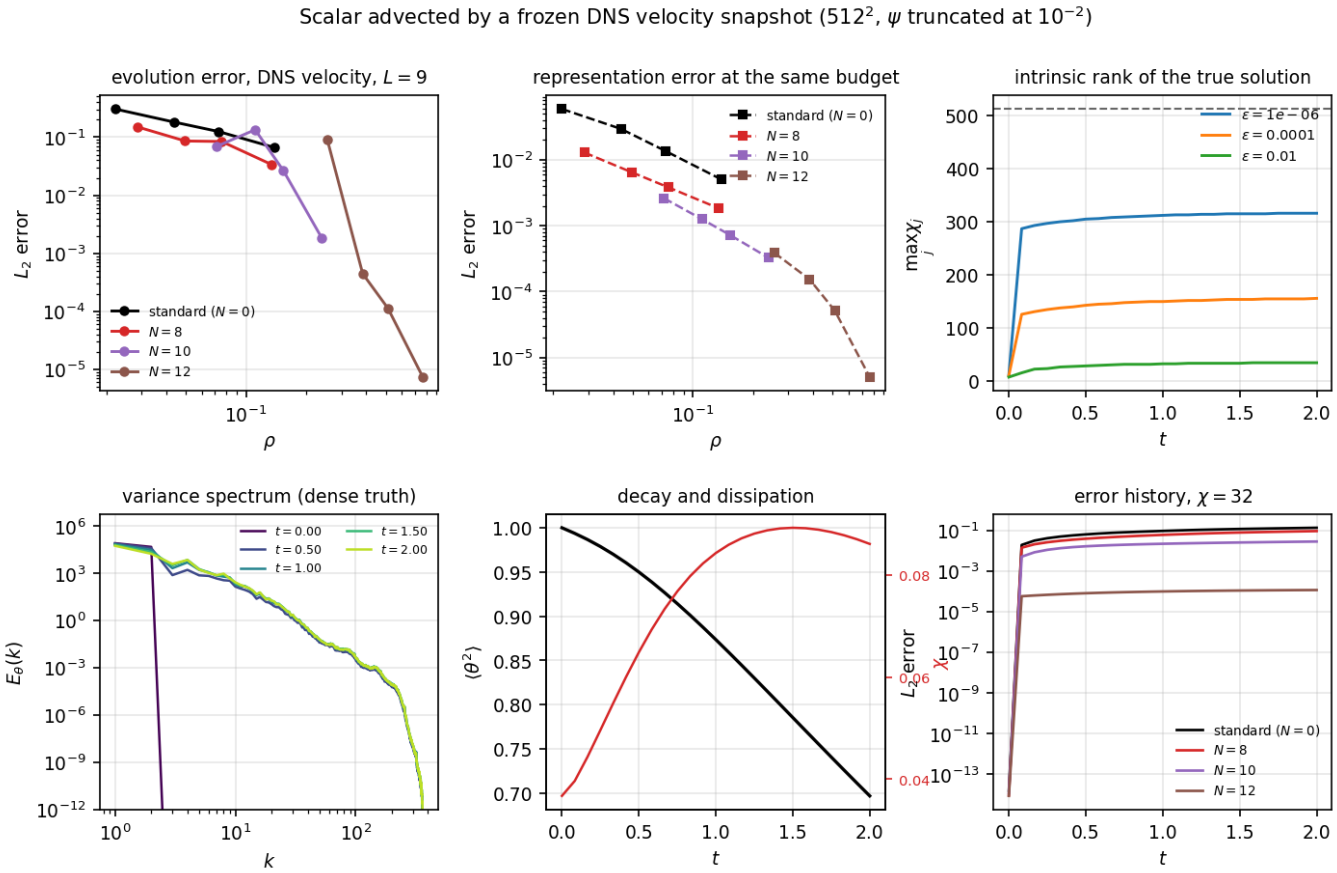


FIG. 4. Scalar advected by a frozen DNS velocity snapshot: rate–distortion, representation error at the same budgets, intrinsic rank, variance spectrum, variance budget and error histories.